


Dockerfiles Best Practices

Nicola Kabar



Agenda

- Introduction to Dockerfiles
- Image Build Process
- Base Images
- Common Dockerfile Directives
- Building Efficient Images

What's a Dockerfile?

- Text document containing the instructions to construct an image
- Docker client sends it along with the context to Docker engine
- Docker engine uses it and a specified *context* to build an image
- Engine goes through a process of running and committing image layers from each line in the Dockerfile
- Uses predefined directives/syntax
- Stored alongside the app's source code in VCS

Sample Dockerfile

```
# A basic apache server. To use either add or bind mount content under /var/www

FROM ubuntu:12.04

MAINTAINER Kimbro Staken version: 0.1

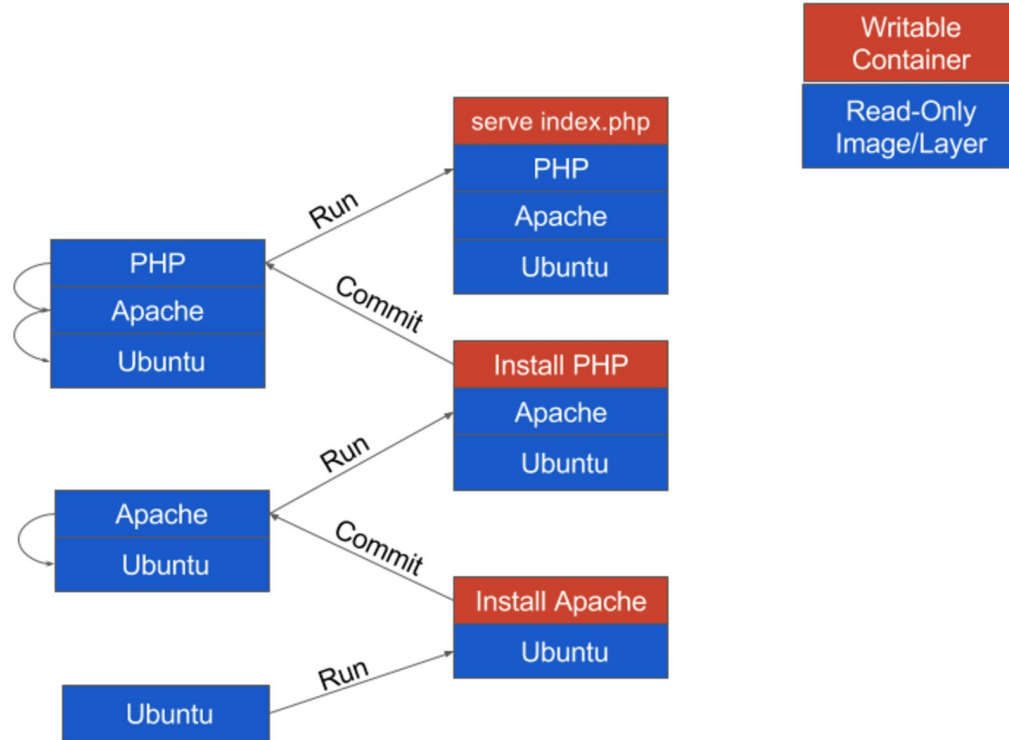

RUN apt-get update && apt-get install -y apache2 && apt-get clean && rm -rf /var/lib/apt/lists/*


ENV APACHE_RUN_USER www-data
ENV APACHE_RUN_GROUP www-data
ENV APACHE_LOG_DIR /var/log/apache2


EXPOSE 80


CMD ["/usr/sbin/apache2", "-D", "FOREGROUND"]
```


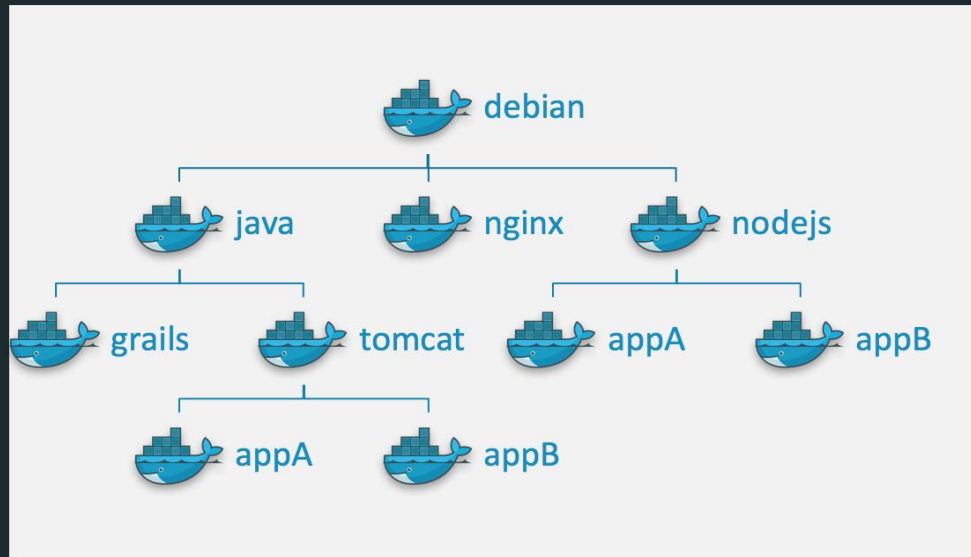
Image Building Process



Base Images

- First line in Dockerfile must always be FROM
- This is the base image in which Docker starts from
- Utilize minimal base images (e.g alpine)
- Recommend to build in-house (middleware, os, systems, ops)
- Should be Scanned and Signed
- Common base images
 - Alpine (all Docker images use alpine)
 - Ubuntu
 - Centos

Hierarchical Order



Common Dockerfile Directives

- 1) **FROM**: specifies base image
- 2) **MAINTAINER**: identifies the maintainer of the image(depreciated :()
- 3) **LABEL**: arbitrary key-value labels for your images
- 4) **RUN**: shell or cmd to run
- 5) **CMD**: instruction to run software of the container
- 6) **ENTRYPOINT**: turns container to an executable
- 7) **EXPOSE**: specify which ports container will listen on
- 8) **ENV**: specify environment variables
- 9) **ADD** or **COPY**: copy files/dirs from context/URL to container
- 10) **USER**: user to use within the container
- 11) **HEALTHCHECK**: custom health-check of container
- 12) **WORKDIR**: working directory within the container
- 13) **SHELL**: specify custom shell (default /bin/sh)
- 14) **ARG**: build-time arguments

Building Efficient Images

1) Minimize Number of Layers

BAD!

```
RUN apt-get update
RUN apt-get install -y wget
RUN rm -rf
/var/lib/apt/lists/*
```

Good!

```
RUN apt-get update &&\
    apt-get install -y wget &&\
    rm -rf /var/lib/apt/lists/*
```

Building Efficient Images

2) Use cache-busting: ensuring you never use cache outdated updates

BAD!

```
RUN apt-get update
RUN apt-get install -y wget
RUN rm -rf
/var/lib/apt/lists/*
```

Good!

```
RUN apt-get update &&\
  apt-get install -y wget &&\
  rm -rf /var/lib/apt/lists/*
```

Building Efficient Images

3) Update by version

Good!

```
RUN apt-get update &&\
    apt-get install -y wget &&\
    rm -rf /var/lib/apt/lists/*
```

Better !

```
RUN apt-get update &&\
    apt-get install -y
wget=1.17.1-1ubuntu1.1 &&\
    rm -rf /var/lib/apt/lists/*
```

Building Efficient Images (cont.)

4) Sort multi-line items to avoid duplication and confusion

```
RUN apt-get update && apt-get install -y \  
    bzip2 \  
    cvs \  
    git \  
    mercurial \  
    subversion
```

Building Efficient Images (cont.)

5) Remove caches and archives during a single RUN command so they are not included in your final image

```
RUN wget -O /tmp/tomcat7.tar.gz  
http://www.us.apache.org/dist/tomcat/tomcat-7/v7.0.63/bin/apache-tomcat-7.0.63.tar.gz &&\  
  cd /opt &&\  
  tar zxf /tmp/tomcat7.tar.gz &&\  
  mv /opt/apache-tomcat* /opt/tomcat && \  
  rm /tmp/tomcat7.tar.gz
```

Building Efficient Images (cont)

6) Separate changes that break the cache to utilize layer caching

BAD!

```
COPY . /usr/src  
RUN npm install
```

Good!

```
COPY package.json /usr/src/package.json  
RUN npm install  
COPY . /usr/src
```

Building Efficient Images (cont)

7) Use a .dockerignore file (supports similar patterns to .gitignore)

```
*.log  
*.git  
...
```

Building Efficient Images (cont)

- 8) Don't boot init because containers model processes not machines.
- 9) Don't upgrade in application image builds. Upgrades should be handled in upstreams base images.
- 10) Use specific tags (prefer git SHAs)
- 11) Use CMD and Entrypoint Together.
- 12) Use non-root user when possible
- 13) One container = one responsibility = one process
- 14) Install what you absolutely **NEED**

References

- 1) docs.docker.com :)
- 2) https://docs.docker.com/engine/userguide/eng-image/dockerfile_best-practices/
- 3) <https://docs.docker.com/engine/reference/builder/#run>
- 4) https://docs.docker.com/engine/userguide/eng-image/dockerfile_best-practices/